

Understanding Random Forests

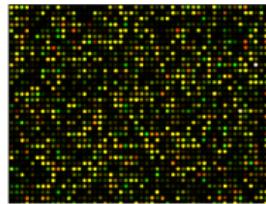
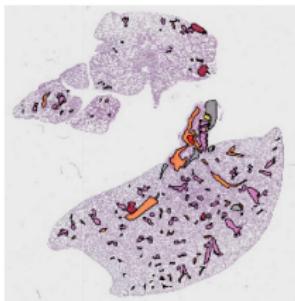
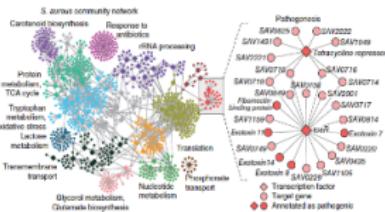
From Theory to Practice

Gilles Louppe

Université de Liège, Belgium

October 9, 2014

Motivation



Objective

From a set of **measurements**,

learn a **model**

to predict and understand **a phenomenon**.

Running example

From **physicochemical properties** (alcohol, acidity, sulphates, ...),

learn a **model**

to predict **wine taste preferences** (from 0 to 10).



P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis, *Modeling wine preferences by data mining from physicochemical properties*, 2009.

Outline

- ① Motivation
- ② Growing decision trees and random forests

Review of state-of-the-art, minor contributions

- ③ Interpreting random forests

Major contributions (Theory)

- ④ Implementing and accelerating random forests

Major contributions (Practice)

- ⑤ Conclusions

Supervised learning

- The **inputs** are random variables $\mathbf{X} = X_1, \dots, X_p$;
- The **output** is a random variable Y .
- Data comes as a finite learning set

$$\mathcal{L} = \{(\mathbf{x}_i, y_i) | i = 0, \dots, N - 1\},$$

where $\mathbf{x}_i \in \mathcal{X} = \mathcal{X}_1 \times \dots \times \mathcal{X}_p$ and $y_i \in \mathcal{Y}$ are randomly drawn from $P_{\mathbf{X}, Y}$.

E.g., $(\mathbf{x}_i, y_i) = ((\text{color} = \text{red}, \text{alcohol} = 12, \dots), \text{score} = 6)$

- The goal is to find a model $\varphi_{\mathcal{L}} : \mathcal{X} \mapsto \mathcal{Y}$ minimizing

$$Err(\varphi_{\mathcal{L}}) = \mathbb{E}_{\mathbf{X}, Y} \{L(Y, \varphi_{\mathcal{L}}(\mathbf{X}))\}.$$

Performance evaluation

Classification

- Symbolic output (e.g., $\mathcal{Y} = \{\text{yes}, \text{no}\}$)
- Zero-one loss

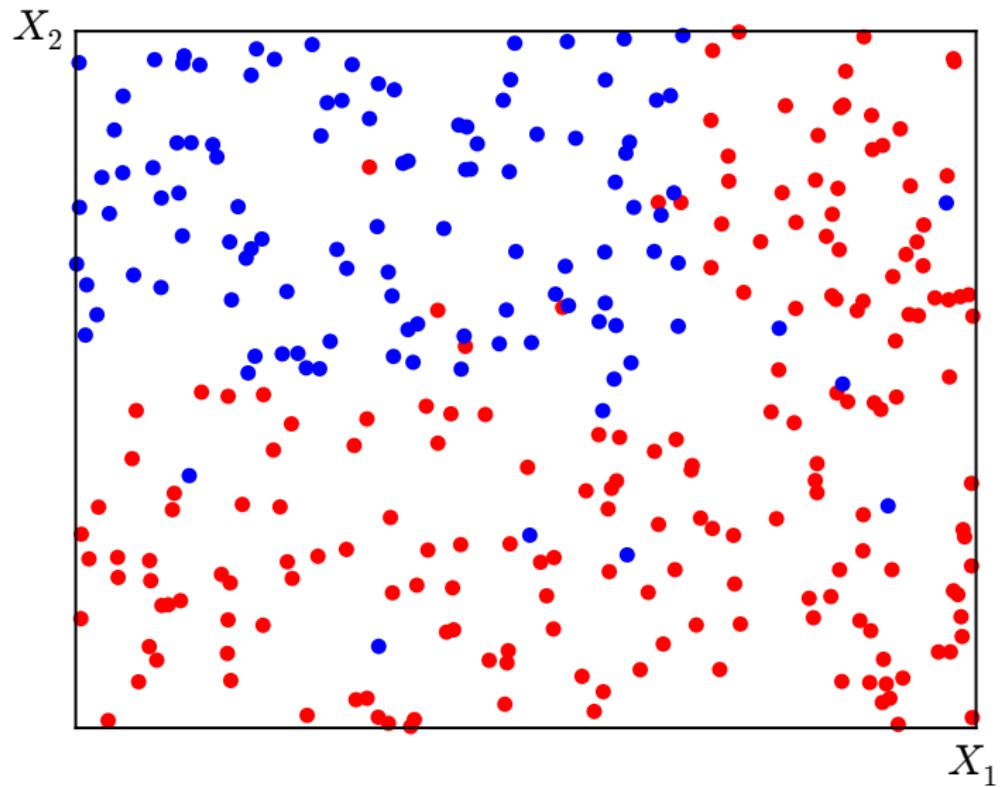
$$L(Y, \varphi_{\mathcal{L}}(X)) = 1(Y \neq \varphi_{\mathcal{L}}(X))$$

Regression

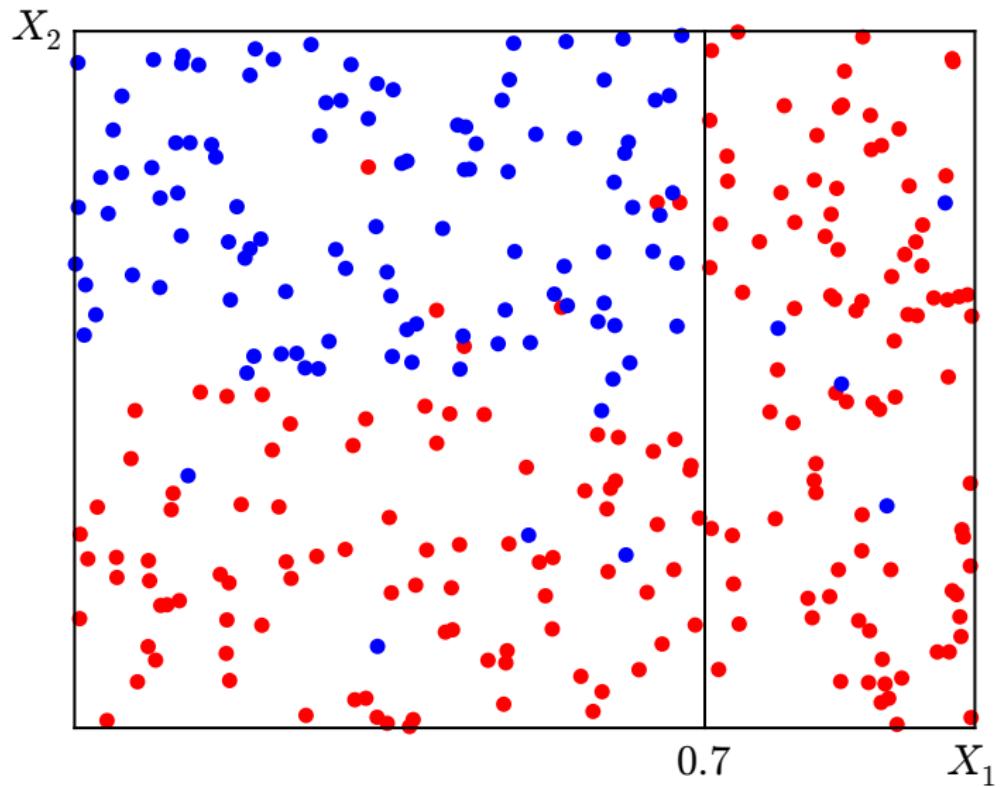
- Numerical output (e.g., $\mathcal{Y} = \mathbb{R}$)
- Squared error loss

$$L(Y, \varphi_{\mathcal{L}}(X)) = (Y - \varphi_{\mathcal{L}}(X))^2$$

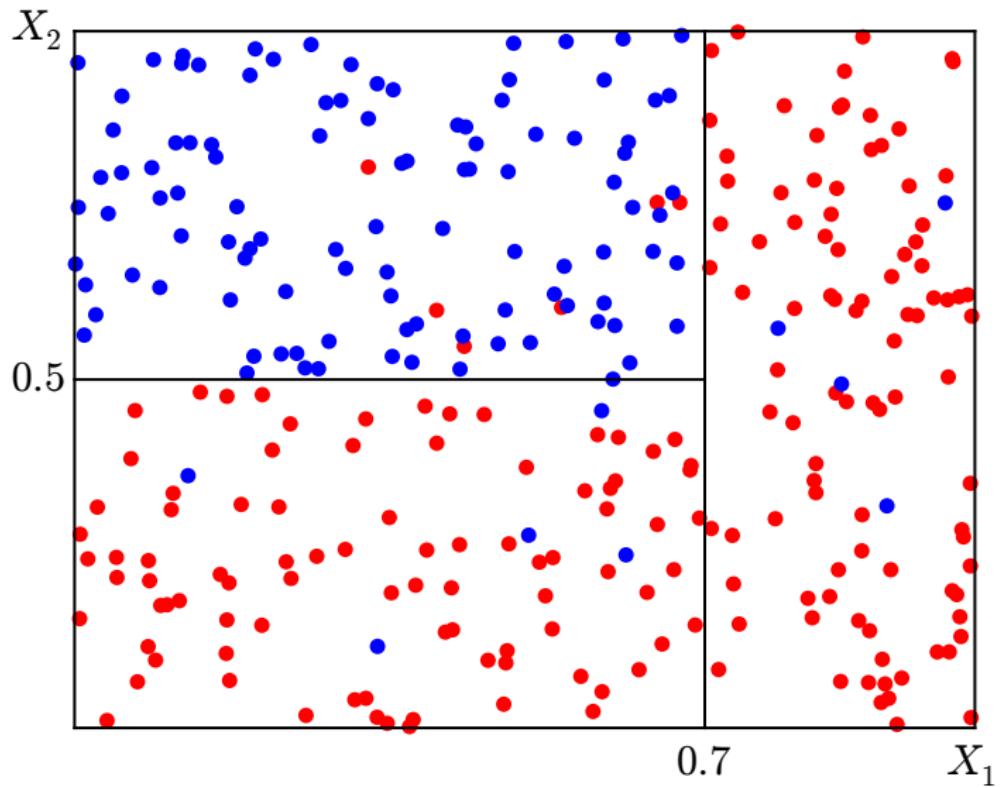
Divide and conquer



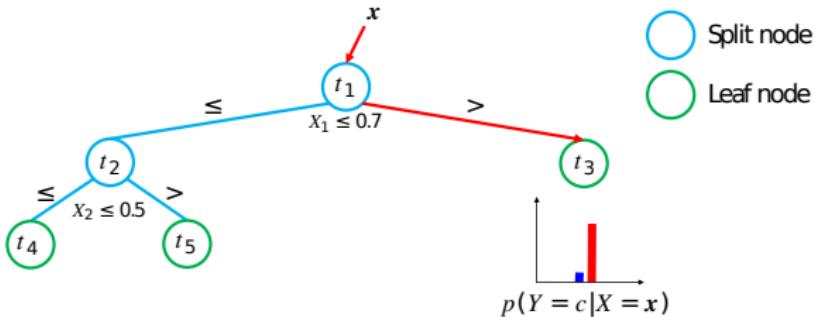
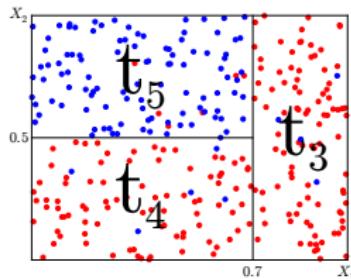
Divide and conquer



Divide and conquer



Decision trees



$t \in \varphi$: nodes of the tree φ

X_t : split variable at t

$v_t \in \mathbb{R}$: split threshold at t

$\varphi(\mathbf{x}) = \arg \max_{c \in \mathcal{Y}} p(Y = c | X = \mathbf{x})$

Learning from data (CART)

function BUILDDECISIONTREE(\mathcal{L})

Create node t from the learning sample $\mathcal{L}_t = \mathcal{L}$

if the stopping criterion is met for t **then**

\hat{y}_t = some constant value

else

 Find the split on \mathcal{L}_t that maximizes impurity decrease

$$s^* = \arg \max_{s \in \Omega} \Delta i(s, t)$$

 Partition \mathcal{L}_t into $\mathcal{L}_{t_L} \cup \mathcal{L}_{t_R}$ according to s^*

t_L = BUILDDECISIONTREE(\mathcal{L}_L)

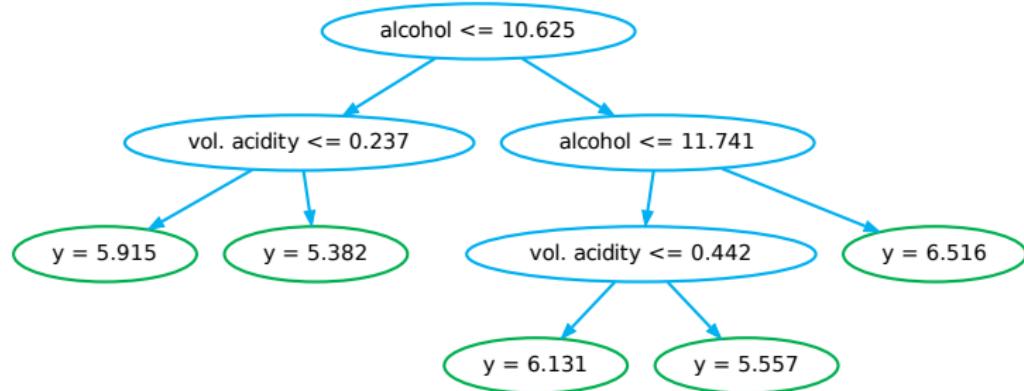
t_R = BUILDDECISIONTREE(\mathcal{L}_R)

end if

return t

end function

Back to our example



Bias-variance decomposition

Theorem. For the squared error loss, the bias-variance decomposition of the expected generalization error at $X = \mathbf{x}$ is

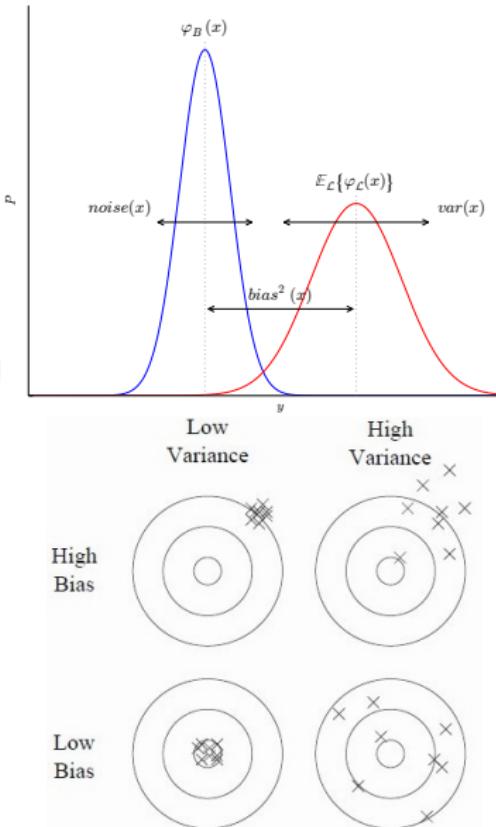
$$\mathbb{E}_{\mathcal{L}}\{Err(\varphi_{\mathcal{L}}(\mathbf{x}))\} = \text{noise}(\mathbf{x}) + \text{bias}^2(\mathbf{x}) + \text{var}(\mathbf{x})$$

where

$$\text{noise}(\mathbf{x}) = Err(\varphi_B(\mathbf{x})),$$

$$\text{bias}^2(\mathbf{x}) = (\varphi_B(\mathbf{x}) - \mathbb{E}_{\mathcal{L}}\{\varphi_{\mathcal{L}}(\mathbf{x})\})^2,$$

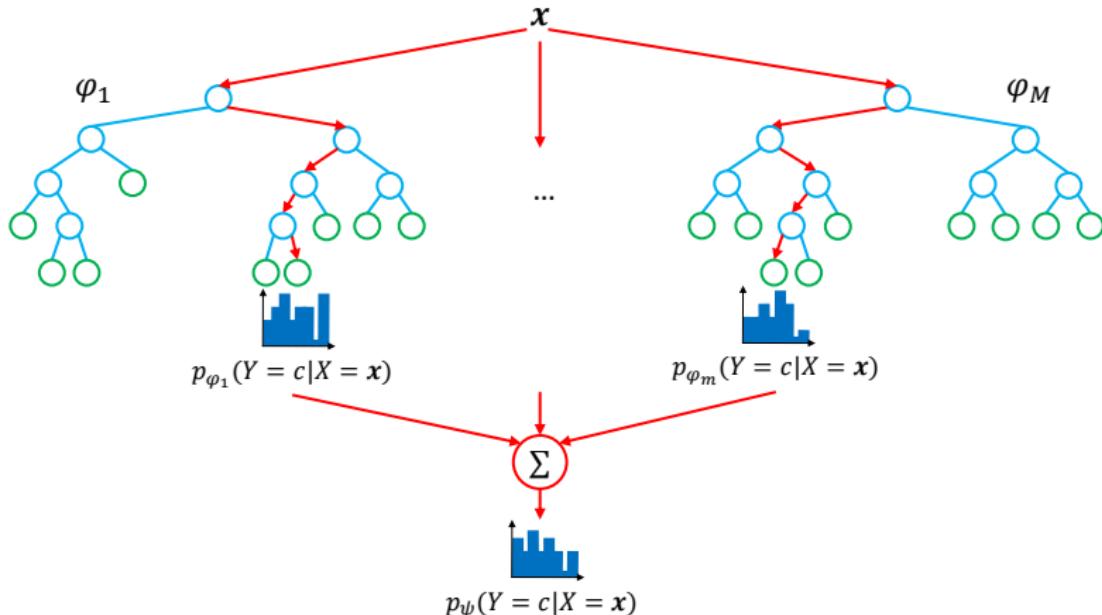
$$\text{var}(\mathbf{x}) = \mathbb{E}_{\mathcal{L}}\{(\mathbb{E}_{\mathcal{L}}\{\varphi_{\mathcal{L}}(\mathbf{x})\} - \varphi_{\mathcal{L}}(\mathbf{x}))^2\}.$$



Diagnosing the generalization error of a decision tree

- (Residual error : Lowest achievable error, independent of $\varphi_{\mathcal{L}}$.)
- Bias : Decision trees usually have **low bias**.
- Variance : They often suffer from **high variance**.
- Solution : *Combine the predictions of several randomized trees into a single model.*

Random forests



Randomization

- Bootstrap samples
- Random selection of $K \leq p$ split variables
- Random selection of the threshold

} Random Forests } Extra-Trees

Bias-variance decomposition (cont.)

Theorem. For the squared error loss, the bias-variance decomposition of the expected generalization error

$\mathbb{E}_{\mathcal{L}}\{Err(\psi_{\mathcal{L}, \theta_1, \dots, \theta_M}(\mathbf{x}))\}$ at $X = \mathbf{x}$ of an ensemble of M randomized models $\varphi_{\mathcal{L}, \theta_m}$ is

$$\mathbb{E}_{\mathcal{L}}\{Err(\psi_{\mathcal{L}, \theta_1, \dots, \theta_M}(\mathbf{x}))\} = \text{noise}(\mathbf{x}) + \text{bias}^2(\mathbf{x}) + \text{var}(\mathbf{x}),$$

where

$$\text{noise}(\mathbf{x}) = Err(\varphi_B(\mathbf{x})),$$

$$\text{bias}^2(\mathbf{x}) = (\varphi_B(\mathbf{x}) - \mathbb{E}_{\mathcal{L}, \theta}\{\varphi_{\mathcal{L}, \theta}(\mathbf{x})\})^2,$$

$$\text{var}(\mathbf{x}) = \rho(\mathbf{x})\sigma_{\mathcal{L}, \theta}^2(\mathbf{x}) + \frac{1 - \rho(\mathbf{x})}{M}\sigma_{\mathcal{L}, \theta}^2(\mathbf{x}).$$

and where $\rho(\mathbf{x})$ is the Pearson correlation coefficient between the predictions of two randomized trees built on the same learning set.

Diagnosing the generalization error of random forests

- Bias : Identical to the bias of a single randomized tree.
- Variance : $\text{var}(\mathbf{x}) = \rho(\mathbf{x})\sigma_{\mathcal{L},\theta}^2(\mathbf{x}) + \frac{1-\rho(\mathbf{x})}{M}\sigma_{\mathcal{L},\theta}^2(\mathbf{x})$
As $M \rightarrow \infty$, $\text{var}(\mathbf{x}) \rightarrow \rho(\mathbf{x})\sigma_{\mathcal{L},\theta}^2(\mathbf{x})$
 - The stronger the randomization, $\rho(\mathbf{x}) \rightarrow 0$, $\text{var}(\mathbf{x}) \rightarrow 0$.
 - The weaker the randomization, $\rho(\mathbf{x}) \rightarrow 1$, $\text{var}(\mathbf{x}) \rightarrow \sigma_{\mathcal{L},\theta}^2(\mathbf{x})$

Bias-variance trade-off. Randomization increases bias but makes it possible to reduce the variance of the corresponding ensemble model. The crux of the problem is to **find the right trade-off**.

Back to our example

<i>Method</i>	<i>Trees</i>	<i>MSE</i>
CART	1	1.055
Random Forest	50	0.517
Extra-Trees	50	0.507

Combining several randomized trees indeed works better !

Outline

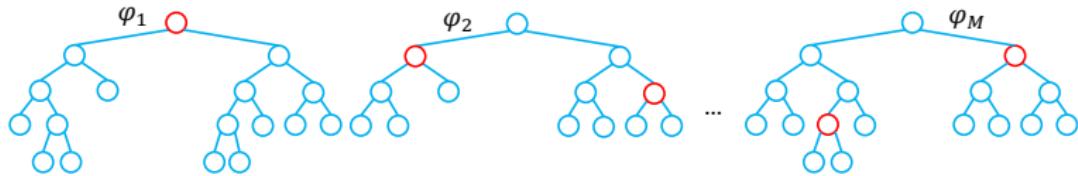
- 1 Motivation
- 2 Growing decision trees and random forests
- 3 Interpreting random forests**
- 4 Implementing and accelerating random forests
- 5 Conclusions

Variable importances



- Interpretability can be recovered through **variable importances**
- Two main importance measures :
 - **The mean decrease of impurity (MDI)** : summing total impurity reductions at all tree nodes where the variable appears (Breiman et al., 1984) ;
 - **The mean decrease of accuracy (MDA)** : measuring accuracy reduction on out-of-bag samples when the values of the variable are randomly permuted (Breiman, 2001).
- We focus here on MDI because :
 - It is faster to compute ;
 - It does not require to use bootstrap sampling ;
 - In practice, it correlates well with the MDA measure.

Mean decrease of impurity



Importance of variable X_j for an ensemble of M trees φ_m is :

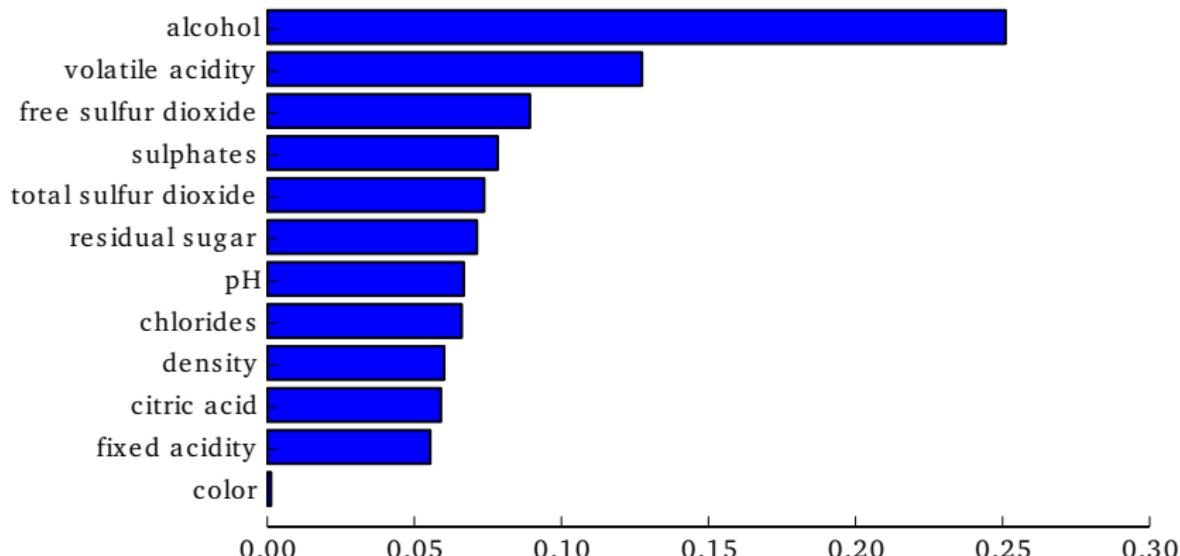
$$\text{Imp}(X_j) = \frac{1}{M} \sum_{m=1}^M \sum_{t \in \varphi_m} 1(j_t = j) [p(t) \Delta i(t)],$$

where j_t denotes the variable used at node t , $p(t) = N_t/N$ and $\Delta i(t)$ is the impurity reduction at node t :

$$\Delta i(t) = i(t) - \frac{N_{t_L}}{N_t} i(t_L) - \frac{N_{t_R}}{N_t} i(t_R)$$

Back to our example

MDI scores as computed from a forest of 1000 fully developed trees on the Wine dataset (Random Forest, default parameters).



What does it mean ?

- MDI works well, but it is not well understood theoretically ;
- We would like to better characterize it and derive its main properties from this characterization.
- Working assumptions :
 - All variables are discrete ;
 - Multi-way splits à la C4.5 (i.e., one branch per value) ;
 - Shannon entropy as impurity measure :

$$i(t) = - \sum_c \frac{N_{t,c}}{N_t} \log \frac{N_{t,c}}{N_t}$$

- **Totally randomized trees (RF with $K = 1$) ;**
- **Asymptotic conditions : $N \rightarrow \infty$, $M \rightarrow \infty$.**

Result 1 : Three-level decomposition (Louppe et al., 2013)

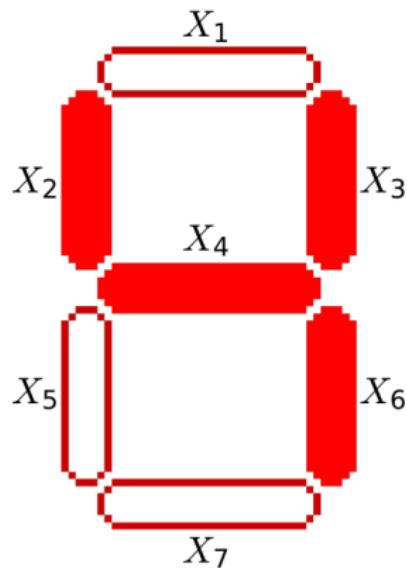
Theorem. Variable importances provide a three-level decomposition of the information jointly provided by all the input variables about the output, accounting for all interaction terms in a fair and exhaustive way.

$$\underbrace{I(X_1, \dots, X_p; Y)}_{\text{Information jointly provided by all input variables about the output}} = \sum_{j=1}^p \underbrace{\text{Imp}(X_j)}_{\text{i) Decomposition in terms of the MDI importance of each input variable}}$$

$$\text{Imp}(X_j) = \sum_{k=0}^{p-1} \underbrace{\frac{1}{C_p^k} \frac{1}{p-k}}_{\text{ii) Decomposition along the degrees } k \text{ of interaction with the other variables}} \underbrace{\sum_{B \in \mathcal{P}_k(V^{-j})} I(X_j; Y|B)}_{\text{iii) Decomposition along all interaction terms } B \text{ of a given degree } k}$$

E.g. : $p = 3$, $\text{Imp}(X_1) = \frac{1}{3}I(X_1; Y) + \frac{1}{6}(I(X_1; Y|X_2) + I(X_1; Y|X_3)) + \frac{1}{3}I(X_1; Y|X_2, X_3)$

Illustration : 7-segment display (Breiman et al., 1984)

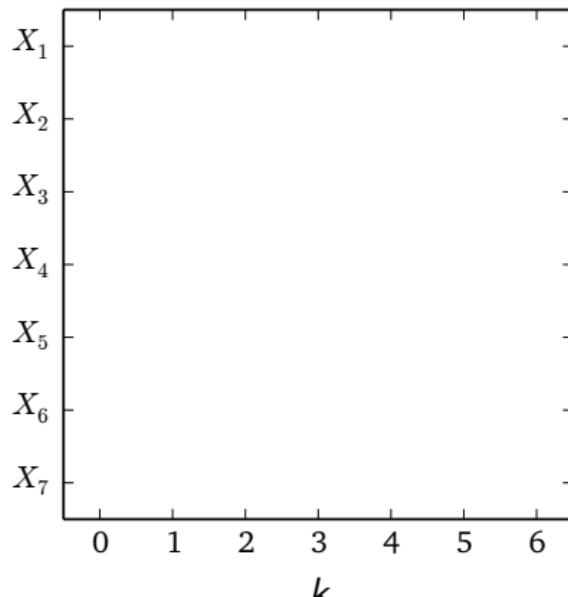


y	x_1	x_2	x_3	x_4	x_5	x_6	x_7
0	1	1	1	0	1	1	1
1	0	0	1	0	0	1	0
2	1	0	1	1	1	0	1
3	1	0	1	1	0	1	1
4	0	1	1	1	0	1	0
5	1	1	0	1	0	1	1
6	1	1	0	1	1	1	1
7	1	0	1	0	0	1	0
8	1	1	1	1	1	1	1
9	1	1	1	1	0	1	1

Illustration : 7-segment display (Breiman et al., 1984)

$$\text{Imp}(X_j) = \sum_{k=0}^{p-1} \frac{1}{C_p^k} \frac{1}{p-k} \sum_{B \in \mathcal{P}_k(V^{-j})} I(X_j; Y|B)$$

Var	Imp
X_1	0.412
X_2	0.581
X_3	0.531
X_4	0.542
X_5	0.656
X_6	0.225
X_7	0.372
\sum	3.321



Result 2 : Irrelevant variables (Louppe et al., 2013)

Theorem. Variable importances depend only on the relevant variables.

Theorem. A variable X_j is irrelevant if and only if $\text{Imp}(X_j) = 0$.

⇒ The importance of a relevant variable is insensitive to the addition or the removal of irrelevant variables.

Definition (Kohavi & John, 1997). A variable X is *irrelevant* (to Y with respect to V) if, for all $B \subseteq V$, $I(X; Y|B) = 0$. A variable is *relevant* if it is not irrelevant.

Relaxing assumptions

When trees are not totally random...

- There can be **relevant variables with zero importances** (due to masking effects).
- The importance of relevant variables can be **influenced by the number of irrelevant variables**.

When the learning set is finite...

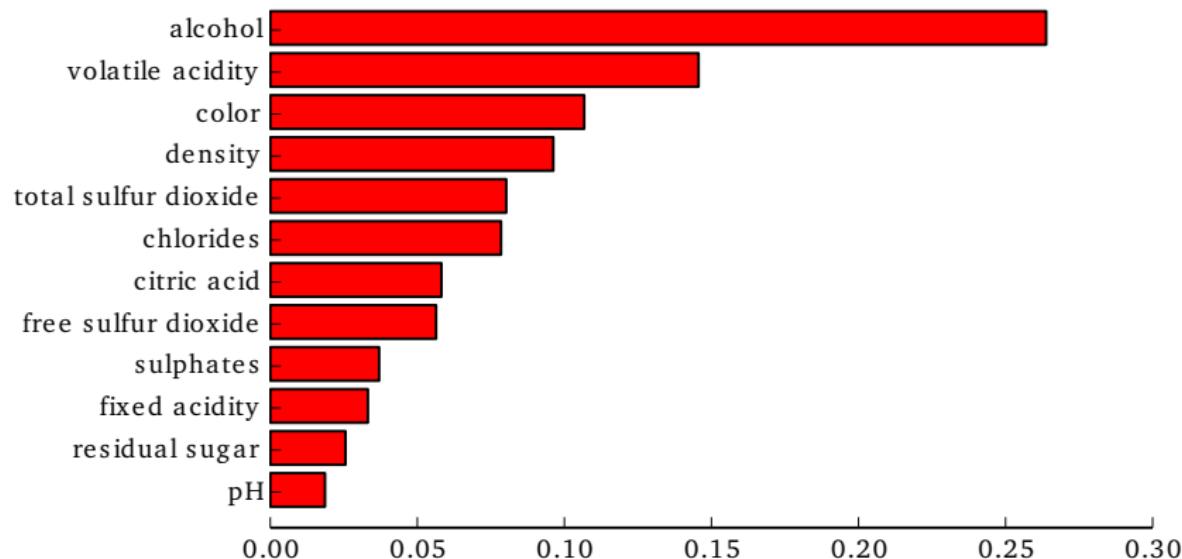
- Importances are **biased towards variables of high cardinality**.
- This effect can be minimized by collecting impurity terms measured from large enough sample only.

When splits are not multiway...

- $i(t)$ does not actually measure the mutual information.

Back to our example

MDI scores as computed from a forest of 1000 fixed-depth trees on the Wine dataset (Extra-Trees, $K = 1$, `max_depth = 5`).



Taking into account (some of) the biases
results in quite a different story !

Outline

- 1 Motivation
- 2 Growing decision trees and random forests
- 3 Interpreting random forests
- 4 Implementing and accelerating random forests
- 5 Conclusions

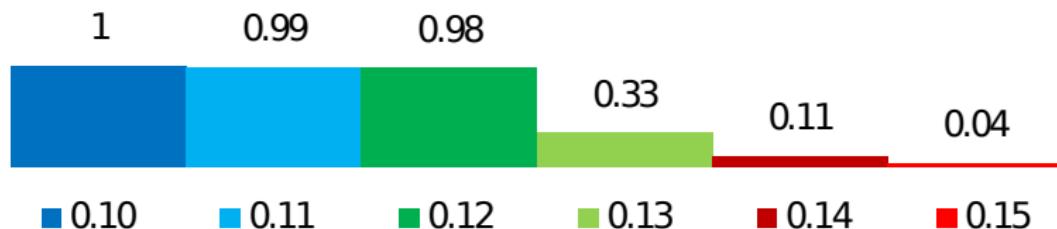


Scikit-Learn

- Open source **machine learning** library for Python
- Classical and **well-established algorithms**
- Emphasis on **code quality** and **usability**

A long team effort

Time for building a Random Forest (relative to version 0.10)



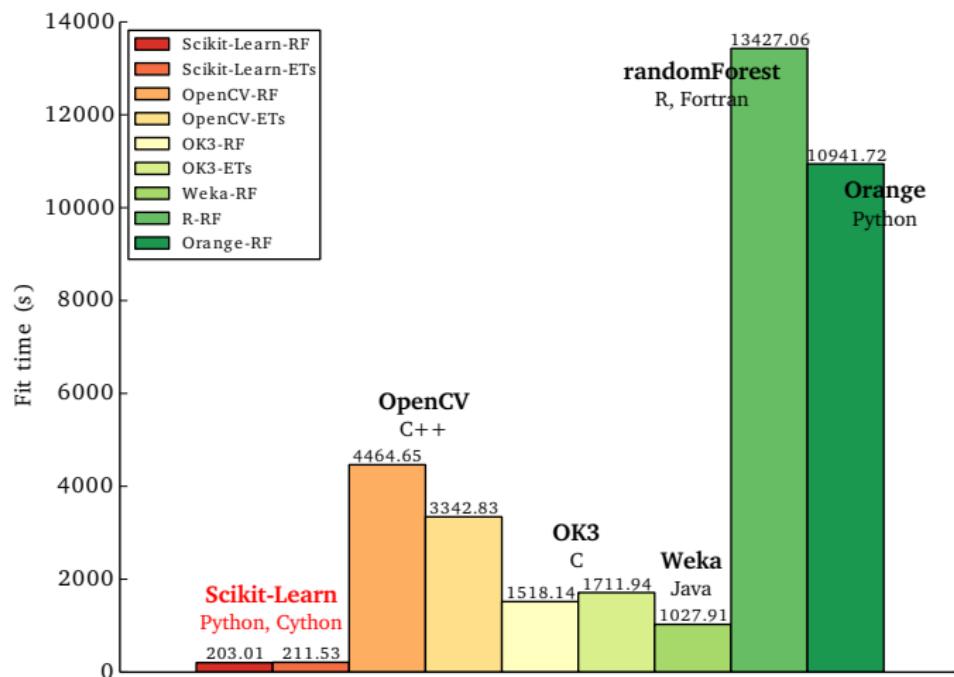
Implementation overview

- Modular implementation, designed with a strict **separation of concerns**
 - Builders : for building and connecting nodes into a tree
 - Splitters : for finding a split
 - Criteria : for evaluating the goodness of a split
 - Tree : dedicated data structure
- Efficient **algorithmic formulation** [See [Louppe, 2014](#)]
 - Dedicated sorting procedure
 - Efficient evaluation of consecutive splits
- **Close to the metal**, carefully coded, implementation
2300+ lines of Python, 3000+ lines of Cython, 1700+ lines of tests

```
# But we kept it stupid simple for users!
clf = RandomForestClassifier()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
```

A winning strategy

Scikit-Learn implementation proves to be **one of the fastest** among all libraries and programming languages.



Computational complexity (Louppe, 2014)

	<i>Average time complexity</i>
CART	$\Theta(pN \log^2 N)$
Random Forest	$\Theta(MK\tilde{N} \log^2 \tilde{N})$
Extra-Trees	$\Theta(MKN \log N)$

- N : number of samples in \mathcal{L}
- p : number of input variables
- K : the number of variables randomly drawn at each node
- $\tilde{N} = 0.632N$.

Improving scalability through randomization

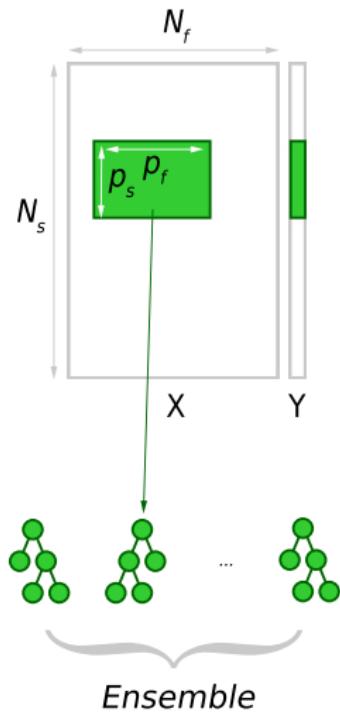
Motivation

- Randomization and averaging allow to improve accuracy by reducing variance.
- As a nice side-effect, the resulting algorithms are fast and embarrassingly parallel.
- Why not purposely exploit randomization to make the algorithm even more scalable (and at least as accurate) ?

Problem

- Let assume a supervised learning problem of N_s samples defined over N_f features. Let also assume T computing nodes, each with a memory capacity limited to M_{max} , with $M_{max} \ll N_s \times N_f$.
- How to best exploit the memory constraint to obtain the most accurate model, as quickly as possible ?

A straightforward solution : Random Patches (Louppe et al., 2012)

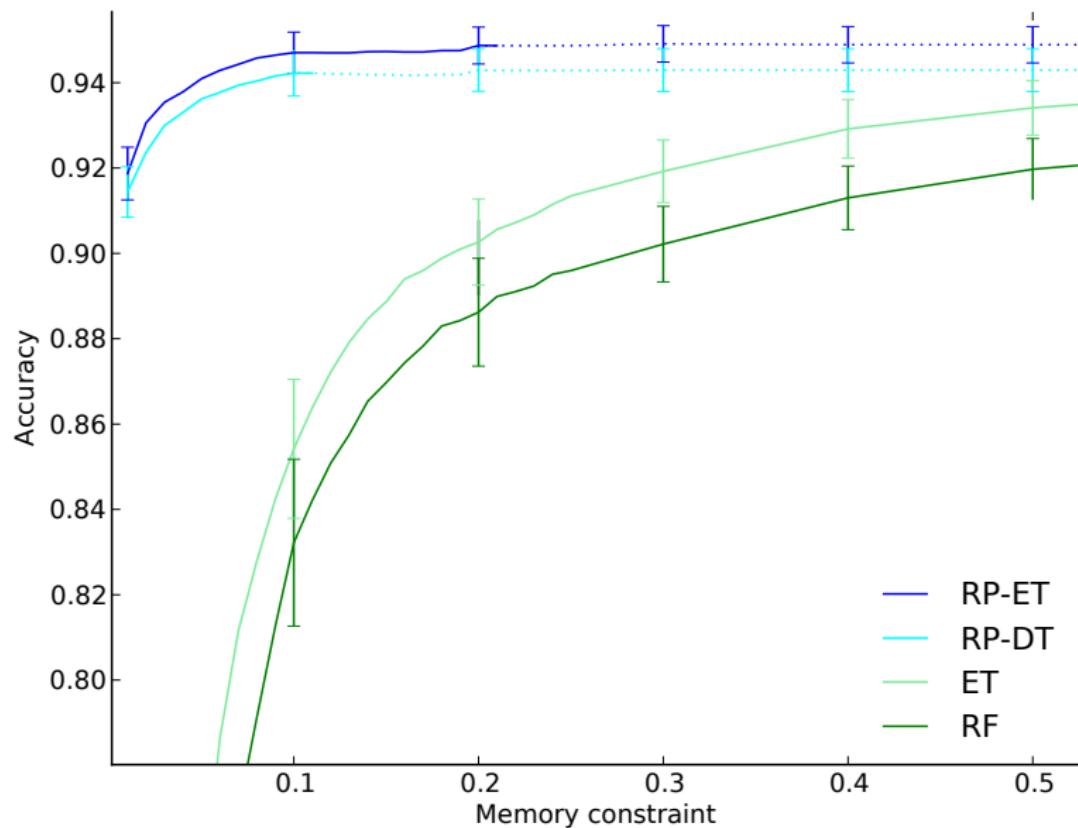


1. Draw a subsample r of $p_s N_s$ random examples, with $p_f N_f$ random features.
2. Build a *base estimator* on r .
3. Repeat 1-2 for a number T of estimators.
4. Aggregate the predictions by voting.

p_s and p_f are two meta-parameters that should be selected

- such that $p_s N_s \times p_f N_f \leq M_{max}$
- to optimize accuracy

Impact of memory constraint



Lessons learned from subsampling

- Training each estimator on the whole data is (often) useless.
The size of the random patches can be reduced without (significant) loss in accuracy.
- As a result, both memory consumption and training time can be reduced, at low cost.
- With strong memory constraints, RP can exploit data better than the other methods.
- Sampling features is critical to improve accuracy. Sampling the examples only is often ineffective.

Outline

- 1 Motivation
- 2 Growing decision trees and random forests
- 3 Interpreting random forests
- 4 Implementing and accelerating random forests
- 5 Conclusions

Opening the black box

- Random forests constitute one of the most **robust and effective** machine learning algorithms for many problems.
- While simple in design and easy to use, random forests remain however
 - hard to analyze theoretically,
 - non-trivial to interpret,
 - difficult to implement properly.
- Through an in-depth re-assessment of the method, this dissertation has proposed **original contributions** on these issues.



Future works

Variable importances

- Theoretical characterization of variable importances in a finite setting.
- (Re-analysis of) empirical studies based on variable importances, in light of the results and conclusions of the thesis.
- Study of variable importances in boosting.

Subsampling

- Finer study of subsampling statistical mechanisms.
- Smart sampling.

Questions ?

Backup slides

Condorcet's jury theorem

Let consider a group of M voters.

If each voter has an independent probability $p > \frac{1}{2}$ of voting for the correct decision, then adding more voters increases the probability of the majority decision to be correct.

When $M \rightarrow \infty$, the probability that the decision taken by the group is correct approaches 1.



Interpretation of $\rho(\mathbf{x})$ (Louppe, 2014)

$$\text{Theorem. } \rho(\mathbf{x}) = \frac{\mathbb{V}_{\mathcal{L}}\{\mathbb{E}_{\theta|\mathcal{L}}\{\varphi_{\mathcal{L},\theta}(\mathbf{x})\}\}}{\mathbb{V}_{\mathcal{L}}\{\mathbb{E}_{\theta|\mathcal{L}}\{\varphi_{\mathcal{L},\theta}(\mathbf{x})\}\} + \mathbb{E}_{\mathcal{L}}\{\mathbb{V}_{\theta|\mathcal{L}}\{\varphi_{\mathcal{L},\theta}(\mathbf{x})\}\}}$$

In other words, it is the ratio between

- the variance due to the learning set and
- the total variance, accounting for random effects due to both the learning set and the random perturbations.

$\rho(\mathbf{x}) \rightarrow 1$ when variance is mostly due to the learning set ;

$\rho(\mathbf{x}) \rightarrow 0$ when variance is mostly due to the random perturbations ;

$\rho(\mathbf{x}) \geq 0$.